

Competitive overview

Lighthouse for AI agents. Open spec, open CLI, public leaderboard.

a14y is an open-source agent-readability scorecard: 38 fast checks across 3 layers, scanning any deployed site for whether AI agents (Claude, Perplexity, ChatGPT) can discover, parse, and act on it. The output is a per-page numeric score with JSON for CI, plus a public 238-site leaderboard. The comparison below is the most current view of where a14y sits relative to other named scorecards in the space.

	a14y	ora.ai (formerly ora.run)	Cloudflare isitagentready.com	agentready.dev	scan-then-sell ¹
Open implementation	✓	✗	✗	✗	✗
Public leaderboard	✓	✓	✗	✗	✗
Per-page scores	✓	✗	✗	✓	✗
CI-friendly (sub-second)	✓	✗	✗	✗	✗
Vendor-neutral	✓	~	✗	~	✗
Sells the fix?	✗	~	~	~	✓

Legend: ✓ yes · ✗ no · ~ partial. ¹ Aggregate of sitespeak.ai, StartDesigns, agnt.one Site Checker, and similar lead-gen scanners — see Wider landscape, page 2.

ORA.AI (FORMERLY ORA.RUN)

The most fully-realized agent-readiness product on the market. Separates an open standard — **AgentReady**, MIT-licensed, vendor-neutral, semver-governed — from a proprietary reference implementation (**Deep Scan v1.1**, 113 checks across 5 weighted layers, F→A+ letter grades, 9,982 sites scanned). Live agent execution (real OAuth, real MCP, multi-turn goal completion across LLM providers) is its wedge; a14y intentionally stays on fast static checks. Recently rebranded with a sharper emphasis on Answer Engine Optimization.

CLOUDFLARE · ISITAGENTREADY.COM

Launched 2026-04-17 during Cloudflare’s Agents Week. Free, proprietary, MCP-queryable; five RFC-aligned dimensions (Discoverability, Content, Bot Access Control, Protocol Discovery, Commerce). No leaderboard — one-shot scans only. Cloudflare positions it as “Lighthouse for AI agents,” the same narrative slot a14y stakes out, but funnels users into Workers / AI Gateway / URL Scanner. Score formula not publicly documented.

AGENTREADY.DEV

Single-purpose 7-check scanner focused on Markdown content negotiation and adjacent signals. Top-of-funnel for **TokenCut**, a commercial prompt-compression proxy. Closer to a single check group of a14y packaged as a stand-alone product than a peer scorecard.

Wider landscape

Seven additional entrants surfaced in the May 2026 sweep. None of them shifts a14y's four moats; together they sharpen them.

Stytch — auth-as-a-service for AI agents (OAuth, MCP, agent detection). Not a scanner; sells the auth layer that ora.ai, Cloudflare, and a14y all **check for**. Featured in Anthropic's agentreadyapps.com masterclass (Episodes 4 and 7).

ambient-code/agentready — MIT open-source scanner with 128 stars on GitHub. Scans **git repositories** for AI-assisted-development readiness (CLAUDE.md, test execution, type annotations) — a fundamentally different axis from a14y's deployed-site scanner. Name-space adjacency, not competition.

sitespeak.ai and **StartDesigns** — scan-then-sell-the-fix lead-gen for a chatbot SaaS and an agency, respectively. Both bundle remediation as the product.

agnt.one Site Checker — diagnostic transparency rather than scoring; shows the raw text payload an agent extracts ("AI Knowledge Preview"). The single strongest idea from the sweep — a14y will likely adopt an equivalent "AI View" mode.

agentreadyapps.com — **not a competitor**. An Anthropic-curated masterclass featuring Stytch, Cloudflare, Arcade, Honeycomb, Langflow, and Stainless. Strategic signal: the agent-readiness ecosystem is being curated, and the open-scorecard slot is currently unclaimed.

What a14y is not chasing

Stated up front so the comparison stays bounded — these are deliberate non-goals, not gaps.

- › **Live OAuth + payment scoring** — ora.ai's territory; not a14y's audience.
- › **Hosted MCP server** — Cloudflare's funnel mechanism; a14y's audience runs the CLI in CI.
- › **Auth-for-agents infrastructure** — Stytch's territory. a14y scores the **presence** of agent-auth, vendor-agnostic.
- › **Remediation-as-a-service** — sitespeak.ai, StartDesigns. a14y's value is vendor-neutral measurement, not selling the fix.
- › **Repository scanning** — ambient-code/agentready owns that axis; a14y scans deployed sites.
- › **A paid backend** — open implementation is the moat.

For press inquiries, source data, or an updated revision:

agentreadability@gmail.com

a14y.dev · github.com/timothyjordan/a14y